



**TRUBA**  
Turkish Science e-Infrastructure



# COMPUTHON 2022/1



ORTA DOĞU TEKNİK ÜNİVERSİTESİ  
MIDDLE EAST TECHNICAL UNIVERSITY



# Computhon Schedule

— — —

**May 4th (Wednesday, Noon):**

Registration deadline

**May 5th (Thursday):**

Hackathon launch (online)

**May 10th (Tuesday, Noon):**

Last day of solution submission (online through Computhon webpage)

**May 11th–12th (Wednesday–Thursday):**

Solution presentations & final evaluation (face-to-face and online)

**May 13th (Friday):**

Announcement of winners & award ceremony (face-to-face and online)

# Today's Program

— — —

- Tensors, fibers, and **problem definition**
- Submission specifications
- Baseline

# Tensors

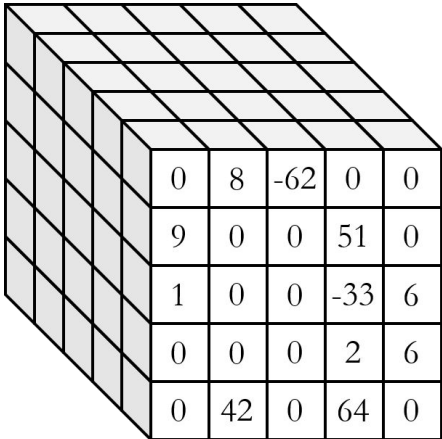
---

**Tensors** are multidimensional arrays, i.e., generalizations of matrices.

- The **order** of a tensor is the number of dimensions (or modes).

0	10	0	0	94	0
-6	0	4	0	0	0
0	0	505	0	0	83
-75	0	0	0	7	0
0	1	0	0	21	11
0	0	0	0	0	0

order-2 tensor



0	8	-62	0	0
9	0	0	51	0
1	0	0	-33	6
0	0	0	2	6
0	42	0	64	0

order-3 tensor

# Tensors

---

Tensors are often stored in-memory in the coordinate (COO) format.

- Coordinate matrix of size (order x #non-zeros)
- The coordinates of each non-zero is a column.
- Each row represents a dimension (mode).
- Values often stored separately (for data type reasons).

0	10	0	0
-6	0	4	0
0	0	5.5	0
0	0	0	0

```
COO =  [0,  1,  1,  2]
       [1,  0,  2,  2]
```


```
vals = [10, -6,  4,  5.5]
```

# Fibers

— — —

**Fibers** are vectors within tensors defined by fixing every dimension's index except one (e.g., rows and columns in matrices).

$T[:,j]$   
(columns)




$T[i,:]$   
(rows)




# Fibers


$T[:,j,z]$   
(columns)



$T[i][:][z]$   
(rows)



$T[i][j][:]$   
(tubes)



# Problem definition

---

- Given a tensor, count the number of fibers that contain at least one non-zero value.

	5	6	7		8	9
1	0	10	0	0	94	0
2	-6	0	4	0	0	0
3	0	0	505	0	0	83
	0	0	0	0	0	0
4	0	1	0	0	21	11
	0	0	0	0	0	0



# Specifications

---

- You are not responsible for reading the tensor files (.tns).
- You will process a tensor stored in the C00 format.
- The values array will not be given with the C00 matrix (and is not needed).

**Expected output: a single integer which is the number of non-zero fibers.**

# Provided data

---

- Main program + baseline fiber counting function.
- Example tensor inputs.
- Sample SLURM scripts for execution.
- OpenMP, MPI, and CUDA training material.

Link: <https://github.com/SU-HPC/Computhon2022-1>

# Categories

— — —

- Single-node: Parallel computation on multi-core CPUs.
- Single-node: Parallel computation on multi-core CPUs and multiple GPUs.
- Multi-node: Distributed computing on multiple nodes + CPUs and GPUs

# Grading criteria

---

- Test tensors have at most 5 billion non-zeros and 10 modes.
- Average speedup on all of our testing tensors (will not be shared with you!)
- Best average speedup wins.
- In case your code times-out/returns incorrect results, you will be penalized (will count as a 0.5x speedup).

# Submission format

---

- For each category:
  - One code and make file.
  - The code file must have **an identical main function** to the one in the given solution file.
    - If you think that you cannot build the solution without modifying the main function, you need to get permission from us first.
  - A README.txt file for each category to give compilation and running details.

# Baseline

— — —

```
/*  
Parameters:  
* [In]  nnz: unsigned int - number of non-zeros  
* [In]  order: unsigned int - number of dimensions  
* [In]  dimensions: unsigned int* - an array of size `order` containing the dimension on each mode  
* [In]  coord: unsigned int** - 2D structure of dimensions `order` x `nnz` with each column  
representing a non-zero  
      ex: coord[i][j] = the i'th dimension value of the j'th non-zero in the tensor  
  
Returns:  
* unsigned long long - number of non-zero fibers in the tensor  
  
*/  
unsigned long long naive_count_coo(unsigned int nnz, unsigned int order, unsigned int* dimensions,  
unsigned int ** coord);
```

```
unsigned long long naive_count_coo(unsigned int nnz, unsigned int order, unsigned int* dimensions, unsigned int **
coord){
    unsigned long long total_fibers = 0;
    // Go over each mode (find fibers on that mode)
    for (int m = 0; m < order; m++){
        // modes other than `m` (dimensions that will be constant for fibers on mode `m`)
        vector<int> modes_to_check;
        for (int x = 0; x < order; x++) {
            if (x == m) continue;
            modes_to_check.push_back(x);
        }
        // Set containing all the fibers on mode `m` based on the constant indices of each fiber
        unordered_set<vector<unsigned int>, vector_hash<unsigned int>, vector_equal<unsigned int>> fibers;
        vector<unsigned int> one_nnz(order-1);
        // Determine the fiber that each non-zero belongs to and add it to the set
        for (unsigned int i = 0; i < nnz; i++){
            // Construct the fiber ID using its dimensions on all modes except `m`
            for (int o = 0; o < order-1; o++){
                one_nnz[o] = coord[modes_to_check[o]][i];
            }
            fibers.insert(one_nnz);
        }
        // Add the fiber count in mode `m` to the total fiber count
        total_fibers+=fibers.size();
    }
    return total_fibers;
}
```